

# Introduction

---

This file captures the notes and important concepts extracted from book Linux Device Driver 3e, essentially acting as a quick cheat sheet Targeted for Linux 6.6 LTS — Hybrid C + Rust Reference

---

## Linux Device Drivers 3rd Ed. — Modern Ready Reckoner (Part 1)

---

**Covers Chapters 1–5 (Blocking I/O, Timers, Memory, Hardware I/O, Interrupts)**

Chapter	Focus
6	Advanced blocking and non-blocking I/O (wait queues, async notifications)
7	Timers and deferred work (bottom halves, tasklets, workqueues)
8	Memory management for drivers (kmalloc, vmalloc, mmap, DMA basics)
9	Hardware I/O and port access
10	Interrupt handling (request_irq, threaded IRQs, shared interrupts)

---

## Chapter 1 — Introduction to Device Drivers

---

Device drivers are the bridge between user space and hardware. They run in kernel space and expose an interface (usually via /dev/\* files) that user applications can call through system calls (open, read, write, ioctl, etc.).

### Key Concepts

- **Kernel vs User Space:** kernel has full access to hardware / memory; user space does not.
- **Driver Types:**
  - **Character Drivers** — byte stream devices (UART, I2C, SPI).
  - **Block Drivers** — storage devices (SD, SATA).
  - **Network Drivers** — packet-oriented.
- **Device Node:** created under /dev; represents interface to the driver.
- **Major / Minor Numbers:** kernel identifies driver and device instance.
- **Modular Kernel:** loadable modules (.ko) can be inserted or removed dynamically.

### Minimal Example (C)

```
// hello.c
#include <linux/module.h>
#include <linux/init.h>
```

```

static int __init hello_init(void) {
    pr_info("Hello, kernel world!\n");
    return 0;
}

static void __exit hello_exit(void) {
    pr_info("Goodbye, kernel world!\n");
}

module_init(hello_init);
module_exit(hello_exit);

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Dhruv");
MODULE_DESCRIPTION("Minimal example module for LDD3 reckoner");

```

Build and load

```

make -C /lib/modules/$(uname -r)/build M=$PWD modules
sudo insmod hello.ko
sudo rmmod hello
dmesg | tail

```

## Rust Version

```

use kernel::prelude::*;

module! {
    type: HelloModule,
    name: b"hello_rust",
    author: b"Dhruv",
    description: b"Hello World in Rust Kernel",
    license: b"GPL",
}

struct HelloModule;

impl KernelModule for HelloModule {
    fn init() -> Result<Self> {
        pr_info!("Hello from Rust kernel world!\n");
        Ok(HelloModule)
    }
}

```

## Learning Notes

- Start every driver journey with a “Hello World” module to verify build → load → unload cycle.
- Understand that each module lives in kernel address space → bugs can crash the system.

- `pr_info()` replaces `printk(KERN_INFO ...)` in newer kernels.
- All kernel code must be GPL-compatible if it exports symbols.

## Real-World Reference:

Look at `/drivers/char/random.c` for an example of a basic char driver implementation.

# Chapter 2 — Building and Loading Modules

Modules are built outside the main kernel tree but linked against its headers using kbuild.

### Key Files

```
# Makefile
obj-m := hello.o
```

Run

```
make -C /lib/modules/$(uname -r)/build M=$PWD modules
```

## Module Management Commands

Command	Action
<code>insmod &lt;module.ko&gt;</code>	Load module manually
<code>modprobe &lt;module&gt;</code>	Load module + dependencies
<code>rmmmod &lt;module&gt;</code>	Remove module
<code>lsmod</code>	List loaded modules
<code>modinfo &lt;module.ko&gt;</code>	View module metadata

## Signing Modules (Linux 6.6)

When Secure Boot is enabled, modules must be signed:

```
scripts/sign-file sha256 key.priv key.x509 hello.ko
```

## Learning Notes

- `modprobe` reads `/lib/modules/{ver}/modules.dep` to resolve dependencies.
- `MODULE_LICENSE("GPL")` ensures kernel exports are accessible.

- Keep builds out-of-tree using the M= syntax to avoid polluting the kernel source.

## Book Examples

Example Name	Brief Description	GitHub Source
hello	Simplest “Hello World” loadable module using <code>printk</code> .	<a href="#">martinezjavier/ldd3 – ch2/hello.c</a>
hello_param	Demonstrates passing parameters via <code>insmod</code> ( <code>module_param</code> ).	<a href="#">ldd3/misc-modules/hellop.c</a>
export / use_export	Shows how one module exports symbols for another to use.	<a href="#">ldd3/misc-modules/export.c</a>

## Real World:

All modern distro kernels use DKMS (Dynamic Kernel Module Support) to rebuild external modules automatically on kernel updates.

---

## Chapter 3 — Character Drivers

Character drivers transfer data as a stream of bytes and are registered with the kernel via `cdev`.

### Key Concepts

- **Device Registration:**

```
static dev_t dev;
alloc_chrdev_region(&dev, 0, 1, "simple_char");
```

- **Character Device Structure:**

```
static struct cdev c_dev;
cdev_init(&c_dev, &fops);
cdev_add(&c_dev, dev, 1);
```

- **File Operations:** Each open device file is associated with a set of callbacks (`file_operations`).

### Example C Driver

```
#include <linux/module.h>
#include <linux/fs.h>
#include <linux/cdev.h>
#include <linux/uaccess.h>
```

```

#define BUF_SIZE 128
static char device_buffer[BUF_SIZE];
static struct cdev my_cdev;
static dev_t dev;

static ssize_t my_read(struct file *filp, char __user *buf, size_t len,
loff_t *off)
{
    size_t bytes = min(len, (size_t)(BUF_SIZE - *off));
    if (copy_to_user(buf, device_buffer + *off, bytes))
        return -EFAULT;
    *off += bytes;
    return bytes;
}

static ssize_t my_write(struct file *filp, const char __user *buf, size_t len,
loff_t *off)
{
    size_t bytes = min(len, (size_t)(BUF_SIZE - *off));
    if (copy_from_user(device_buffer + *off, buf, bytes))
        return -EFAULT;
    *off += bytes;
    return bytes;
}

static struct file_operations fops = {
    .owner = THIS_MODULE,
    .read = my_read,
    .write = my_write,
};

static int __init my_init(void)
{
    alloc_chrdev_region(&dev, 0, 1, "simple_char");
    cdev_init(&my_cdev, &fops);
    cdev_add(&my_cdev, dev, 1);
    pr_info("Registered char dev %d:%d\n", MAJOR(dev), MINOR(dev));
    return 0;
}

static void __exit my_exit(void)
{
    cdev_del(&my_cdev);
    unregister_chrdev_region(dev, 1);
    pr_info("simple_char removed\n");
}

module_init(my_init);
module_exit(my_exit);
MODULE_LICENSE("GPL");

```

Create the device node:

```
sudo mknod /dev/simple_char c <major> 0
```

Test

```
echo "abc" | sudo tee /dev/simple_char
sudo cat /dev/simple_char
```

## Learning Notes

- `copy_to_user` / `copy_from_user` handle address space protection.
- Always check user-supplied lengths and pointers.
- Each char driver instance typically corresponds to one minor number.

## Book Examples

Example Name	Brief Description	GitHub Source
<code>scull</code>	"Simple Character Utility for Loading Localities" — core character driver with <code>open</code> , <code>read</code> , <code>write</code> , etc.	<a href="#">ldd3/scull/</a>
<code>scullp</code> , <code>scullc</code> , <code>sculld</code> , <code>scullv</code>	Variants of <code>scull</code> showing different memory management models (page-based, chunked, linked, virtual).	<a href="#">ldd3/scull/</a>

## Rust Kernel Analogy

```
use kernel::file::FileOperations;

struct MyChar;
impl FileOperations for MyChar {
    kernel::declare_file_operations!(read, write);
    fn read(...) -> Result<usize> { /* similar logic */ }
}
```

---

## Chapter 4 — Debugging Techniques

When developing kernel drivers, bugs can crash the entire system — so you need safe methods to trace behavior, inspect memory, and log intelligently.

### Logging and Tracing

Use the kernel's built-in logging functions (preferred over `printf`)

```
pr_debug("debug: x=%d\n", x);
pr_info("info: device opened\n");
pr_warn("warning: low buffer\n");
pr_err("error: invalid state\n");
```

Enable dynamic debugging for specific modules:

```
echo 'file simple_char.c +p' > /sys/kernel/debug/dynamic_debug/control
```

To disable:

```
echo 'file simple_char.c -p' > /sys/kernel/debug/dynamic_debug/control
```

## Using dmesg

All pr\_\* output goes to the kernel ring buffer, viewable via:

```
dmesg | tail
```

To clear:

```
sudo dmesg -C
```

## Kernel Oops and Backtraces

If your driver dereferences an invalid pointer or uses an uninitialized structure, you'll get an Oops message. A sample backtrace:

```
BUG: unable to handle kernel NULL pointer dereference at 00000000
IP: [<fffffffff8123b2e0>] my_read+0x14/0x80 [simple_char]
...

```

Use addr2line to map it back:

```
addr2line -e simple_char.ko 0x14
```

## Kernel Probes (kprobes, ftrace, bpf)

For modern kernels, you can dynamically instrument functions.

### Example: ftrace

```
echo function > /sys/kernel/debug/tracing/current_tracer
echo my_read > /sys/kernel/debug/tracing/set_ftrace_filter
cat /sys/kernel/debug/tracing/trace
```

### Rust Example (for inline tracing)

```
kernel::pr_debug!("MyChar::read() called");
```

## Debugging Tools Summary

Tool	Use Case
dmesg, pr_info()	Basic kernel logs
dynamic_debug	Enable/disable runtime debug logs
ftrace	Function call tracing
kgdb	Step-through debugging
perf	Performance profiling
bpftrace	Dynamic event tracing

## Learning Notes

- Never use printf or std::println() in kernel space — only kernel-safe logging macros.
- Keep pr\_debug lines in code but disable them via config — useful later in production.
- If your driver crashes, check /var/log/kern.log or journalctl -k.
- For hardware I/O debugging, hexdump, strace, and logic analyzers are invaluable.

## Book Examples

Example Name	Brief Description	GitHub Source
faulty	Module that deliberately crashes to demonstrate kernel debugging.	<a href="#">mharsch/ldd3-samples/faulty</a>
oops	Demonstrates triggering oops and inspecting kernel backtraces.	<a href="#">jesstess/ldd3-examples/faulty</a>

## Real-World Reference:

Look at `drivers/tty/serial/serial_core.c` — a masterclass in debug logging and safe instrumentation.

---

## Chapter 5 — File Operations Deep Dive

---

A driver connects with user space primarily through the `file_operations` structure. This structure defines how the kernel calls your driver on user actions (open, read, write, ioctl, poll, mmap).

### Common File Operations

Operation	Description	User API Trigger
<code>.open</code>	Initialize or prepare device	<code>open()</code>
<code>.release</code>	Cleanup	<code>close()</code>
<code>.read</code>	Transfer data to user	<code>read()</code>
<code>.write</code>	Receive data from user	<code>write()</code>
<code>.unlocked_ioctl</code>	Custom control commands	<code>ioctl()</code>
<code>.poll</code>	Event-based readiness	<code>select(), poll()</code>
<code>.mmap</code>	Map device memory to user space	<code>mmap()</code>

### Example: ioctl Handling

```
#define IOCTL_CLEAR _IO('a', 1)
#define IOCTL_FILL _IOW('a', 2, int)

static long my_ioctl(struct file *f, unsigned int cmd, unsigned long arg)
{
    switch (cmd) {
    case IOCTL_CLEAR:
        memset(device_buffer, 0, BUF_SIZE);
        break;
    case IOCTL_FILL: {
        int val;
        if (copy_from_user(&val, (int __user *)arg, sizeof(val)))
            return -EFAULT;
        memset(device_buffer, val, BUF_SIZE);
        break;
    }
    default:
        return -EINVAL;
    }
    return 0;
}

static const struct file_operations fops = {
```

```

    .owner          = THIS_MODULE,
    .read           = my_read,
    .write          = my_write,
    .unlocked_ioctl = my_ioctl,
};


```

Test from user space:

```

int fd = open("/dev/simple_char", O_RDWR);
int val = 0x41;
ioctl(fd, IOCTL_FILL, &val);


```

## Example: Non-blocking I/O & Poll

If you want your driver to support select() or poll(), you can use a wait queue.

```

DECLARE_WAIT_QUEUE_HEAD(wq);
static int flag = 0;

static ssize_t my_read(...) {
    wait_event_interruptible(wq, flag != 0);
    flag = 0;
    return 0;
}

static __poll_t my_poll(struct file *filp, poll_table *wait) {
    poll_wait(filp, &wq, wait);
    if (flag)
        return EPOLLIN | EPOLLRDNORM;
    return 0;
}

```

From user space:

```

poll() waits until data is ready for reading


```

## Learning Notes

- Always validate ioctl command numbers and argument pointers.
- Use `_IOC_DIR`, `_IOC_TYPE`, `_IOC_NR`, `_IOC_SIZE` macros to decode requests.
- Prefer `unlocked_ioctl` over legacy `ioctl` in modern kernels.
- For asynchronous event-driven design, integrate wait queues and poll methods.
- In Rust kernel drivers, these are implemented via traits under `kernel::file`.

## Book Examples

Example Name	Brief Description	GitHub Source
scullconcurrent	Shows safe concurrent access using semaphores and spinlocks.	<a href="#">ldd3/scull/</a>
scullpipe	Blocking I/O driver using wait queues for read/write synchronization.	<a href="#">ldd3/scull/pipe.c</a>

## Real-World Reference

- `/drivers/input/evdev.c` — perfect example of `poll`, `read`, and `ioctl`.
- `/drivers/tty/tty_io.c` — full implementation of file ops for terminal devices.

---

## Summary of Part 1

---

Chapter	Focus	Key Concepts
1	Intro to Drivers	Kernel vs user space, module basics
2	Module Mechanics	Building, inserting, removing
3	Char Drivers	<code>cdev</code> , <code>file_operations</code> , buffers
4	Debugging	<code>pr_debug</code> , ftrace, dynamic_debug
5	File Operations	ioctl, poll, non-blocking I/O

## Learning Map

1. Start with “Hello World” kernel module
2. Learn build/load cycle (`insmod`, `modprobe`)
3. Write a simple char driver (`read`, `write`)
4. Add `ioctl` & `poll`
5. Integrate debugging macros

---

## Part 2 — Linux Device Drivers Ready Reckoner

---

Covers Chapters 6–10 (Blocking I/O, Timers, Memory, Hardware I/O, Interrupts)

Chapter	Focus
6	Advanced blocking and non-blocking I/O (wait queues, async notifications)
7	Timers and deferred work (bottom halves, tasklets, workqueues)
8	Memory management for drivers ( <code>kmalloc</code> , <code>vmalloc</code> , <code>mmap</code> , DMA basics)

## Chapter Focus

---

9	Hardware I/O and port access
10	Interrupt handling (request_irq, threaded IRQs, shared interrupts)

## Targeted for Linux 6.6 LTS — Hybrid C + Rust Reference

# Chapter 6 — Blocking and Non-blocking I/O

---

Device drivers often need to suspend a process until a condition is met, e.g., data available for reading.

## Wait Queues

Wait queues let a process sleep until a condition becomes true.

```
DECLARE_WAIT_QUEUE_HEAD(my_waitqueue);
static int flag = 0;

static ssize_t my_read(struct file *filp, char __user *buf, size_t len,
loff_t *off)
{
    wait_event_interruptible(my_waitqueue, flag != 0);
    flag = 0;
    return 0;
}
```

- `wait_event_interruptible(queue, condition)` sleeps until `condition` is true.
- Non-blocking behavior: return `-EAGAIN` if `O_NONBLOCK` is set.

## Poll and Select

```
static __poll_t my_poll(struct file *filp, poll_table *wait)
{
    poll_wait(filp, &my_waitqueue, wait);
    if (flag)
        return EPOLLIN | EPOLLRDNORM;
    return 0;
}
```

User-space test:

```
fd_set fds;
FD_ZERO(&fds);
FD_SET(fd, &fds);
select(fd+1, &fds, NULL, NULL, NULL);
```

## Learning Notes

- Always guard shared variables with proper synchronization (spinlocks/mutexes).
- Non-blocking I/O is critical for GUI apps or network drivers.

## Book Examples

Example Name	Brief Description	GitHub Source
sculluid	Adds file ownership and user-based access control.	<a href="#">ldd3/scull/sculluid.c</a>
scull_access	Demonstrates file permission control and <code>open</code> policies.	<a href="#">ldd3/scull/scull_access.c</a>
scullsingle	Restricts driver to single/multiple open semantics.	<a href="#">ldd3/scull/</a>

## Chapter 7 — Timers and Deferred Work

Kernel allows deferring work to be done later:

- Timers: Schedule function at specific time.
- Tasklets: Run in softirq context, non-blocking, quick execution.
- Workqueues: Kernel thread context, can sleep.

### Timer Example

```
#include <linux/timer.h>
static struct timer_list my_timer;

void timer_callback(struct timer_list *t)
{
    pr_info("Timer fired!\n");
}

static int __init my_init(void)
{
    timer_setup(&my_timer, timer_callback, 0);
    mod_timer(&my_timer, jiffies + msecs_to_jiffies(2000));
    return 0;
}
```

### Tasklet Example

```
#include <linux/interrupt.h>

void tasklet_func(unsigned long data) {
    pr_info("Tasklet executed: %lu\n", data);
}
```

```

}

DECLARE_TASKLET(my_tasklet, tasklet_func, 42);
tasklet_schedule(&my_tasklet);

```

## Workqueue Example

```

#include <linux/workqueue.h>

static void my_work_func(struct work_struct *work) {
    pr_info("Workqueue executed\n");
}
static DECLARE_WORK(my_work, my_work_func);
schedule_work(&my_work);

```

## Learning Notes

- **Tasklets:** Cannot sleep, run in interrupt context.
- **Workqueues:** Can sleep, run in process context.
- Timers are suitable for delayed actions; tasklets/workqueues for bottom halves.

## Book Examples

Example Name	Brief Description	GitHub Source
scullm	Demonstrates mapping kernel memory to user space with <code>mmap</code> .	<a href="#">ldd3/scull/mmap.c</a>

## Chapter 8 — Memory Management

Drivers allocate memory for buffers, device structures, and DMA.

### kmalloc / kfree

```

char *buf = kmalloc(128, GFP_KERNEL);
if (!buf) return -ENOMEM;
kfree(buf);

```

- `GFP_KERNEL` may sleep, suitable for process context.
- `GFP_ATOMIC` must be used in interrupt context.

### Lookaside Cache (kmem\_cache)

```

struct kmem_cache *my_cache;
my_cache = kmem_cache_create("my_cache", sizeof(struct my_obj), 0,
SLAB_HWCACHE_ALIGN, NULL);
void *obj = kmem_cache_alloc(my_cache, GFP_KERNEL);
kmem_cache_free(my_cache, obj);

```

## vmalloc

Allocates contiguous virtual memory, may not be physically contiguous:

```

void *vbuf = vmalloc(4096);
vfree(vbuf);

```

## Per-CPU Variables

```

DEFINE_PER_CPU(int, my_counter);
this_cpu_inc(my_counter);

```

## Learning Notes

- Always free memory in the same context you allocated.
- DMA buffers may need `dma_alloc_coherent`.
- Avoid `vmalloc` in high-performance, frequently accessed paths.

## Book Examples

<b>Example Name</b>	<b>Brief Description</b>	<b>GitHub Source</b>
<code>kalloc</code>	Demonstrates <code>kmalloc</code> , <code>vmalloc</code> , and slab allocator use.	<a href="#">mharsch/ldd3-samples/memory</a>

## Chapter 9 — Hardware I/O

Device drivers communicate through I/O ports or memory-mapped I/O.

### I/O Port Access

```

#include <asm/io.h>
outb(0xFF, 0x378); // write byte to parallel port
unsigned char val = inb(0x378); // read byte

```

Check allocation:

```
request_region(0x378, 3, "my_port");
release_region(0x378, 3);
```

## Memory-Mapped I/O

```
#include <linux/io.h>
void __iomem *reg;
reg = ioremap(0xFE000000, 0x100);
iowrite32(0x1234, reg);
u32 val = ioread32(reg);
iounmap(reg);
```

## Learning Notes

- Use memory barriers (wmb(), rmb()) to prevent reordering.
- Always request/release resources to avoid conflicts.
- For ISA / legacy ports, check /proc/ioports.

## Book Examples

Example Name	Brief Description	GitHub Source
short	Simulated hardware I/O driver using the parallel port for I/O access and interrupt simulation.	<a href="#">ldd3/short/</a>
shortprint	Debug printing version of <code>short</code> for port I/O visibility.	<a href="#">ldd3/short/shortprint.c</a>

## Chapter 10 — Interrupt Handling

Interrupts allow devices to signal the CPU asynchronously.

### Basic IRQ Handling

```
#include <linux/interrupt.h>

static irqreturn_t my_irq_handler(int irq, void *dev_id)
{
    pr_info("IRQ %d triggered\n", irq);
    return IRQ_HANDLED;
}
```

```
request_irq(17, my_irq_handler, IRQF_SHARED, "my_irq", &my_dev);
free_irq(17, &my_dev);
```

- **Top-half:** quick execution, non-blocking.
- **Bottom-half:** deferred work (tasklets, workqueues).

## Shared IRQs

- Must specify `IRQF_SHARED`.
- Use unique `dev_id`.
- Return `IRQ_NONE` if the interrupt was not for your device.

## Learning Notes

- Avoid sleeping in top-half (use `GFP_ATOMIC` if allocating memory).
- Use tasklets/workqueues for longer processing.
- Use procfs or sysfs to monitor interrupts (`/proc/interrupts`).

## Book Examples

Example Name	Brief Description	GitHub Source
<code>shortirq</code>	Demonstrates interrupt handling and shared IRQs.	<a href="#">ldd3/short/shortirq.c</a>
<code>shortprintirq</code>	Interrupt handler with detailed event logging.	<a href="#">ldd3/short/shortprintirq.c</a>

## Summary of Part 2

Chapter	Focus	Key Points
6	Blocking I/O	Wait queues, poll/select, non-blocking I/O
7	Timers & Deferred Work	Timers, tasklets, workqueues, bottom halves
8	Memory Management	<code>kmalloc</code> , <code>vmalloc</code> , mempools, per-CPU variables
9	Hardware I/O	Ports, memory-mapped I/O, barriers
10	Interrupts	<code>request_irq</code> , top/bottom halves, shared IRQs

## Part 3 — Linux Device Drivers Ready Reckoner

Chapters 11–15 (Kernel Types, Data Structures, USB & Serial Drivers, Synchronization)

Chapter	Focus
11	Kernel Data Types & Endianness

Chapter	Focus
12	Linked Lists & Data Structures
13	USB Drivers & Endpoints
14	I2C, SPI, UART Drivers (Character Devices)
15	Concurrency & Synchronization (Spinlocks, Semaphores, RCU)

## Chapter 11 — Kernel Data Types & Endianness

Kernel provides explicitly-sized data types for portability.

### Standard Types

```
u8 a; // unsigned 8-bit
u16 b; // unsigned 16-bit
u32 c; // unsigned 32-bit
u64 d; // unsigned 64-bit
```

- Defined in `<linux/types.h>`
- `size_t` and `ssize_t` for memory/IO sizes

### Endianness Macros

```
#include <asm/byteorder.h>
u32 val = cpu_to_le32(0x12345678);
u32 v2 = le32_to_cpu(val);
```

### Pointers and Error Values

- Many kernel functions return `ERR_PTR(-errno)` instead of `NULL` for richer error reporting.
- Use `IS_ERR()` and `PTR_ERR()` to handle.

### Learning Notes

- Always prefer explicitly sized types in drivers for portability.
- Be careful with cross-architecture data transfer (USB, network, storage).
- Avoid assumptions about pointer size — use `unsigned long` where necessary.

---

## Chapter 12 — Linked Lists & Data Structures

Linux provides doubly-linked circular lists in `<linux/list.h>`:

```
struct my_node {
    int data;
    struct list_head list;
};

LIST_HEAD(my_list);
```

## Operations

```
struct my_node *n = kmalloc(sizeof(*n), GFP_KERNEL);
n->data = 42;
list_add(&n->list, &my_list);

struct my_node *tmp;
list_for_each_entry(tmp, &my_list, list) {
    pr_info("Value: %d\n", tmp->data);
}
```

## Learning Notes

- List operations are **not thread-safe**; use spinlocks if needed.
- `list_for_each_entry_safe` allows deletion during iteration.

# Chapter 13 — USB Drivers & Endpoints

USB devices communicate through endpoints grouped into interfaces.

## USB Types

- CONTROL — setup & configuration
- BULK — large data transfer (storage devices)
- INTERRUPT — small, periodic (keyboards)
- ISOCHRONOUS — audio/video streaming

## Example USB Skeleton

```
static int my_probe(struct usb_interface *intf, const struct usb_device_id
*id) {
    pr_info("USB device connected\n");
    return 0;
}
```

```

static void my_disconnect(struct usb_interface *intf) {
    pr_info("USB device disconnected\n");
}

static struct usb_driver my_driver = {
    .name = "my_usb",
    .id_table = my_id_table,
    .probe = my_probe,
    .disconnect = my_disconnect,
};

module_usb_driver(my_driver);

```

## USB Request Blocks (URBs)

- Allocate with `usb_alloc_urb()`
- Submit with `usb_submit_urb()`
- Callback on completion for async transfers

## Learning Notes

- Use `usb_control_msg()` for simple requests.
- URBs are essential for async bulk/interrupt transfers.
- Sysfs exposes USB device info (`/sys/bus/usb/devices`).

## Chapter 14 — I2C, SPI, UART Drivers (Character Devices)

These are specialized char drivers interfacing with serial buses.

### I2C Example

```

struct i2c_client *client;
i2c_smbus_write_byte_data(client, reg, value);
int val = i2c_smbus_read_byte_data(client, reg);

```

### SPI Example

```

struct spi_device *spi;
u8 tx[4] = {0x01, 0x02, 0x03, 0x04};
struct spi_transfer t = {
    .tx_buf = tx,
    .len = sizeof(tx),
};
spi_sync_transfer(spi, &t, 1);

```

## UART Example

- Often exposed via `tty_driver`
- Use `struct uart_port` + `struct uart_driver`
- Supports blocking/non-blocking read/write

## Learning Notes

- These drivers are char-device-based and fit into `file_operations`.
- Concurrency control is critical: use mutexes or spinlocks.
- I2C/SPI have **master vs slave** roles — always check bus protocol.
- Rust kernel traits (`i2c::I2cClient`, `spi::SpiDevice`) are analogous.

---

# Chapter 15 — Concurrency & Synchronization

---

Kernel drivers must handle multiple processes and interrupts safely.

## Spinlocks

```
spinlock_t lock;
spin_lock(&lock);
// critical section
spin_unlock(&lock);
```

## Mutexes

```
struct mutex mtx;
mutex_lock(&mtx);
// safe to sleep
mutex_unlock(&mtx);
```

## Read-Copy-Update (RCU)

```
struct my_node *ptr;
rcu_read_lock();
ptr = rcu_dereference(my_head);
rcu_read_unlock();
```

## Learning Notes

- **Interrupt context:** spinlocks + GFP\_ATOMIC
  - **Process context:** mutexes allowed
  - Avoid deadlocks by locking order and minimizing critical section.
  - RCU allows read-mostly data structures without blocking readers.
- 

## Summary of Part 3

---

Chapter	Focus	Key Points
11	Kernel Types	u8/u16/u32/u64, endianness, error pointers
12	Linked Lists	list_head, safe iteration, synchronization
13	USB	Endpoints, URBs, interfaces, async transfers
14	I2C/SPI/UART	Serial bus drivers as char devices
15	Concurrency	Spinlocks, mutexes, RCU, interrupt safety

## Learning Map

1. Use explicit types and handle endianness.
  2. Master linked lists and memory-safe iteration.
  3. Understand USB endpoint communication and URBs.
  4. Serial bus drivers (I2C/SPI/UART) are specialized char drivers.
  5. Apply proper locking primitives based on context (spinlock vs mutex).
- 

## Part 4 — Linux Device Drivers Ready Reckoner

---

Chapters 16–20 (DMA, Memory-Mapped Devices, Power Management, Hotplug, Advanced Topics)

## Chapter 16 — Direct Memory Access (DMA)

---

DMA allows **devices to transfer data directly to/from memory** without CPU intervention.

### DMA Types

- **Consistent / Coherent DMA:** CPU and device always see same data.
- **Streaming / Normal DMA:** CPU must flush caches.

### DMA API Example

### Scatter-Gather DMA

- Supports non-contiguous memory
- Often used in network/storage drivers

- Handled via `struct scatterlist` and `dma_map_sg()`

## Learning Notes

- DMA-safe memory must be allocated with `dma_alloc_coherent` or `dma_map_single`.
- CPU must not directly access device memory mapped for DMA without proper barriers.
- Watch out for cache coherency on non-coherent architectures.

---

# Chapter 17 — Memory-Mapped Devices (MMIO)

---

Many devices expose registers in memory space.

## Mapping MMIO

```
#include <linux/io.h>

void __iomem *regs;
regs = ioremap(phys_addr, size);
iowrite32(0x1234, regs);
u32 val = ioread32(regs);
iounmap(regs);
```

## Learning Notes

- Always use `ioremap` and `iounmap`.
- Use `wmb()` / `rmb()` to enforce write/read ordering.
- MMIO is preferred over legacy port-mapped I/O on modern platforms.

---

# Chapter 18 — Power Management (PM)

---

Drivers must respect system sleep, suspend, and resume.

## PM Callbacks

```
static int my_suspend(struct device *dev) { /* save state */ return 0; }
static int my_resume(struct device *dev) { /* restore state */ return 0; }

static const struct dev_pm_ops my_pm_ops = {
    .suspend = my_suspend,
    .resume   = my_resume,
};
```

## Runtime PM

- Device can autosuspend when idle
- APIs: `pm_runtime_enable()`, `pm_runtime_get_sync()`, `pm_runtime_put_sync()`

## Learning Notes

- Suspend/resume must save device state (registers, buffers).
- Always match runtime PM calls to avoid reference leaks.
- Critical for battery-powered embedded systems.

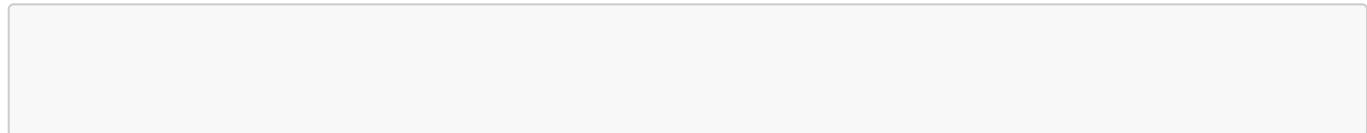
---

## Chapter 19 — Hotplug & Device Registration

---

Linux supports dynamic **plug-and-play** for USB, PCI, I2C, SPI.

### PCI Driver Example



## Learning Notes

- Always match `.probe` and `.remove` lifecycle.
- Use `devm_*` managed resources to simplify cleanup.
- Hotplug devices require careful interrupt and memory management.

---

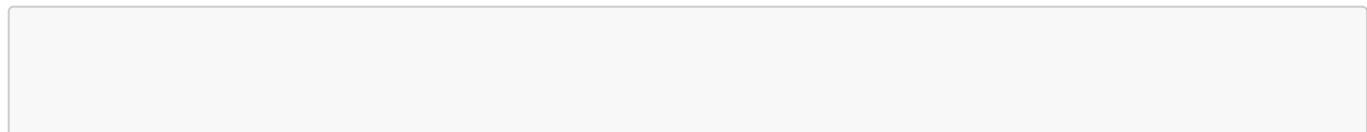
## Chapter 20 — Advanced Topics & Kernel Integration

---

### DMA + Interrupts

- Combine **bottom halves** and DMA for efficient transfers.
- Avoid sleeping in interrupt context; schedule workqueues for heavy processing.

### Memory Barriers & Atomic Operations



### Kernel Modules Best Practices

- Always free memory and IRQs in `exit` function.
- Avoid busy waiting; prefer wait queues or tasklets.
- Test on multiple architectures if using explicit sizes or endianness.

## Learning Notes

- Modern kernel drivers use hybrid `interrupt-driven + DMA` for efficiency.
  - Use `devm_*` APIs to reduce boilerplate and resource leaks.
  - Keep critical sections minimal; prefer async processing for high-speed devices.
- 

## Summary of Part 4

Chapter	Focus	Key Points
16	DMA	<code>dma_alloc_coherent</code> , scatter-gather, CPU/device synchronization
17	MMIO	<code>ioremap</code> , <code>ioread/iowrite</code> , memory barriers
18	Power Management	<code>suspend/resume</code> , runtime PM, autosuspend
19	Hotplug	PCI, USB, device probe/remove, devm resources
20	Advanced Topics	DMA+interrupts, atomic operations, kernel best practices

## Learning Map

1. Understand DMA flows and cache coherency.
  2. Use MMIO safely with proper ordering.
  3. Implement suspend/resume for device power management.
  4. Register devices properly for hotplug support.
  5. Combine interrupts, DMA, and deferred work efficiently.
- 

## Part 5 — LDD3 Mental Map Cheat Sheet

Linux Device Drivers — Kernel Developer Quick Reference Target: Linux 6.6 LTS | Hybrid C/Rust perspective

### 1. Kernel Data & Types

Concept	C Example	Notes / Rust Analogy
Fixed-size integers	<code>u8, u16, u32, u64</code>	Rust: <code>u8/u16/u32/u64</code>
Endianness	<code>cpu_to_le32(val)</code>	Rust: <code>to_le() / from_le()</code>
Error pointers	<code>ERR_PTR(-ENOMEM)</code>	Rust: <code>Result&lt;T, Error&gt;</code>

### 2. Memory Management

Concept	Example	Notes
<code>kmalloc</code>	<code>kmalloc(128, GFP_KERNEL)</code>	GFP_ATOMIC in IRQ, GFP_KERNEL in process context

Concept	Example	Notes
vmalloc	vmalloc(4096)	Virtually contiguous memory
Lookaside cache	kmem_cache_alloc()	High-volume object allocation
Per-CPU vars	DEFINE_PER_CPU(int, counter)	Each CPU has separate copy
DMA-safe	dma_alloc_coherent()	CPU/device coherent memory

### 3. I/O & Character Devices

Bus/Device	Kernel API	Notes
UART	tty_driver, uart_port	Blocking/non-blocking read/write
I2C	i2c_smbus_read/write	Master/slave roles, char-device style
SPI	spi_sync_transfer()	Blocking or async transfers
Generic char	file_operations	.read, .write, .ioctl

### 4. Blocking & Async I/O

Concept	C Example	Notes
Wait queues	wait_event_interruptible(queue, cond)	Sleep until condition
Poll/select	.poll() with poll_wait()	Non-blocking user-space support
Async notification	fasync_struct, kill_fasync()	Signals user-space process

### 5. Timers & Deferred Work

Concept	Example	Notes
Kernel timers	timer_setup() + mod_timer()	Schedule function later
Tasklets	DECLARE_TASKLET(my_tasklet, func, data)	Softirq context, cannot sleep
Workqueues	DECLARE_WORK(work, func)	Process context, can sleep

### 6. Interrupts

Concept	Example	Notes
Top-half	request_irq() handler	Quick response, no sleep
Bottom-half	Tasklets/Workqueues	Deferred heavy processing

Concept	Example	Notes
Shared IRQ	<code>IRQF_SHARED</code> + unique <code>dev_id</code>	Check if interrupt is for your device

## 7. USB Drivers

Concept	Example	Notes
Endpoints	CONTROL/BULK/INT/ISO	Direction: IN/OUT
URB	<code>usb_alloc_urb()</code> + <code>usb_submit_urb()</code>	Async transfer with callback
Interfaces	<code>struct usb_interface</code>	Group endpoints by function

## 8. PCI & Hotplug

Concept	Example	Notes
PCI driver	<code>pci_driver</code> , <code>.probe</code> , <code>.remove</code>	Devm-managed resources simplify cleanup
Hotplug	USB/PCI dynamic detection	Handle attach/detach cleanly

## 9. Power Management

Concept	Example	Notes
Suspend/Resume	<code>.suspend = my_suspend</code>	Save device state
Runtime PM	<code>pm_runtime_get_sync()</code> / <code>pm_runtime_put_sync()</code>	Autosuspend idle devices

## 10. Concurrency & Synchronization

Concept	Example	Notes
Spinlocks	<code>spin_lock() / spin_unlock()</code>	IRQ-safe, cannot sleep
Mutex	<code>mutex_lock() / unlock()</code>	Sleep allowed, process context
RCU	<code>rcu_read_lock() / rcu_read_unlock()</code>	Read-mostly, wait-free reads
Atomic	<code>atomic_inc() / atomic_dec()</code>	For counters in IRQ or SMP

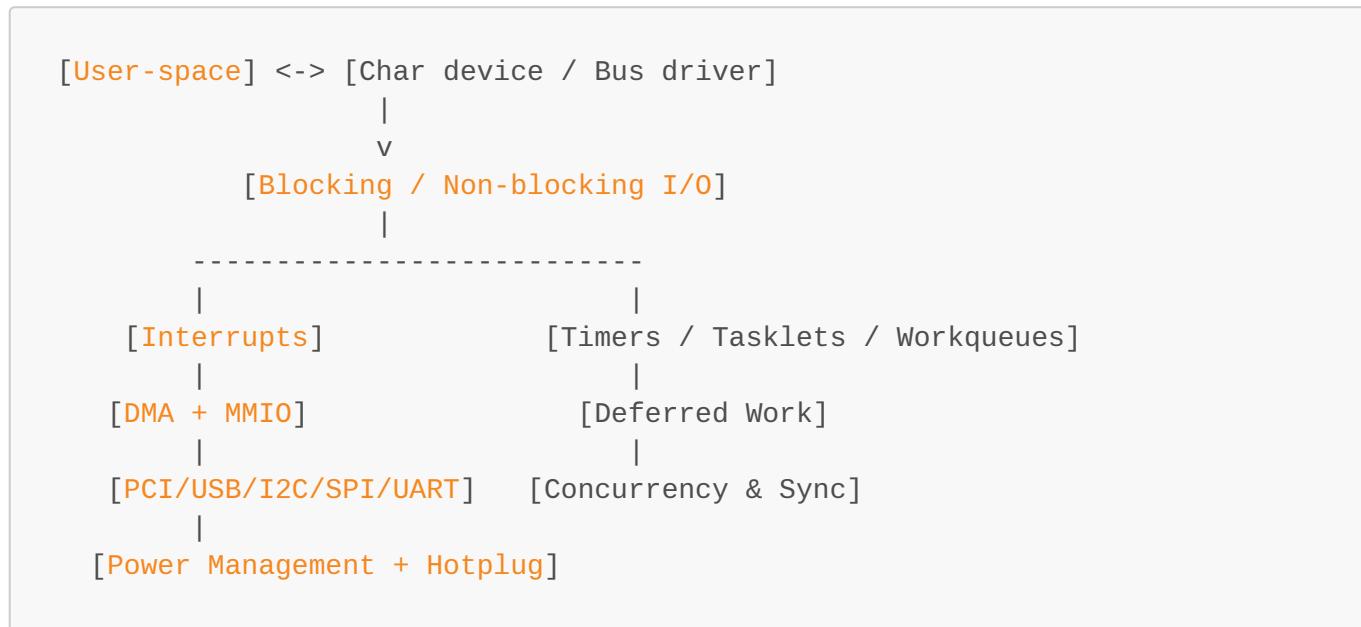
## 11. Advanced DMA + MMIO

Concept	Example	Notes
DMA coherent	<code>dma_alloc_coherent()</code>	CPU & device always see same data
MMIO	<code>ioremap() / iowrite32() / ioread32()</code>	Use memory barriers
Scatter-gather	<code>dma_map_sg()</code>	Non-contiguous memory support

## 12. Best Practices

- Always `release resources` in `exit` or `.remove`.
- Use `devm_` APIs to simplify cleanup.
- Prefer `async + bottom-half` processing over busy-waiting.
- Keep critical sections minimal.
- Test with `non-blocking` and `interrupt contexts`.

## Quick Mental Map Flow



## Rust Kernel Analogies

- `kmalloc` → `alloc::<T>()` / `Box<T>`
- `ioremap` → `IoMem::map()`
- `tasklet` → `KernelTasklet::schedule()`
- `workqueue` → `WorkQueue::queue()`
- `spinlock` → `SpinLock<T>`
- `per-CPU variable` → `PerCpu<T>`